

RECOMMENDATION SYSTEMS

- Imdadullah Khan

Recommendation Systems

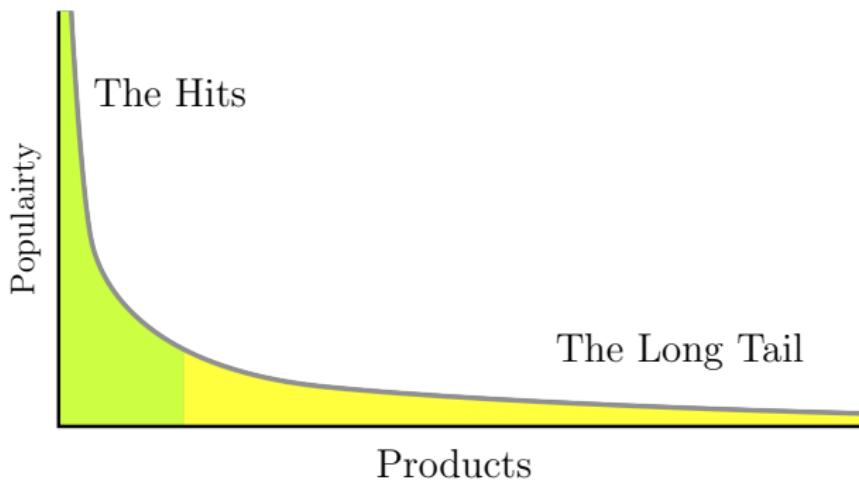


The Web, they say, is leaving the era of search and entering one of discovery. What's the difference? **Search** is what you do when you're looking for something. **Discovery** is when something wonderful that you didn't know existed, or didn't know how to ask for, finds you.

J. O'Brien, Nov 20, 2006 The race to create a 'smart' Google

Recommendation Systems

- Retailers cannot shelve everything
- Online retailers and digital content providers have millions of products



Recommendation Systems

- Near zero-cost dissemination of information about products
- More choice necessitates better information filtering (customization)

Customization can be

- **Hand-Curated:** Chef's specials, editor's picks, favorites
- **Simple aggregates:** Top 10, Trending, Recent uploads
- **Customized to individual users:** Recommendation Systems

Recommendation Systems



2017

Perhaps the single most important algorithmic distinction between “born digital” enterprises and legacy companies is not their people, data sets, or computational resources, but a clear real-time commitment to delivering accurate, actionable customer recommendations.

- **Netflix:** 75% of movies watched are recommended ¹
 - “... personalization and recommendations save us more \$1B per year” ²
- **Amazon:** 35% of purchases on Amazon come from recommendations ¹
- **Google News:** recommendation generate 38% more click-throughs ¹
- **Airbnb:** “Together, Search Ranking and Similar Listings drive 99% of our booking conversions”³
- **Alibaba:** For 11.11. mega sale, targeted personalized landing pages, resulted in 20% higher conversion rate from previous year ⁴

¹ X. Amatriain, (2014) Machine Learning Summer School, CMU

² Gomez-Uribe & Hunt, Netflix Inc., (ACM Trans. on MIS 2015)

³ Grbovic et.al [Airbnb Engineering & Data Science] (2018)

⁴ InsideRetail.Asia (2017)

Recommendation Systems: Problem Formulation

- n users - $\{c_1, \dots, c_n\}$ and m items - $\{p_1, \dots, p_m\}$
- Utility Matrix U : $n \times m$ matrix row/column for each user/item
- $U(i, j)$: rating of user i for item j

	p_1	p_2	p_3	p_j				p_m				
u_1	1		2	1	4		2	3	2	5		2
u_2		1			2	1		2		1		3
u_3		1	1	2		1				1		2
			3	2		5		2		3	4	
	1		2						5			
u_i		3	2	1	4	5	1	3	1	2		1
		4								4		
		5		1						5		
	1		4				1	3		5	1	2
u_n		3		1	1	2	1			4		5

$U(i, j)$ could be

- 0 – 5 stars
- $\in [0, 1]$
- $\in \{0, 1\}$

Computational linear algebra problem of **matrix completion**

Recommendation Systems: Problem Formulation

- n users - $\{c_1, \dots, c_n\}$ and m items - $\{p_1, \dots, p_m\}$
- Utility Matrix U : $n \times m$ matrix row/column for each user/item
- $U(i, j)$: rating of user i for item j

	p_1	p_2	p_3	p_j				p_m				
u_1	1		2	1	4		2	3	2	5		2
u_2		1			2	1		2		1		3
u_3		1	1	2		1				1		2
				3	2		5		2		3	4
		1		2						5		
u_i		3	2	1	4	5	?	1	3	1	2	1
		4								4		
		5		1						5		
	1		4				1	3		5	1	2
u_n		3		1	1	2	1			4		5

$U(i, j)$ could be

- 0 – 5 stars
- $\in [0, 1]$
- $\in \{0, 1\}$

if prediction for $U(i, j)$ is high recommend product j to user i

Recommendation Systems: Challenges

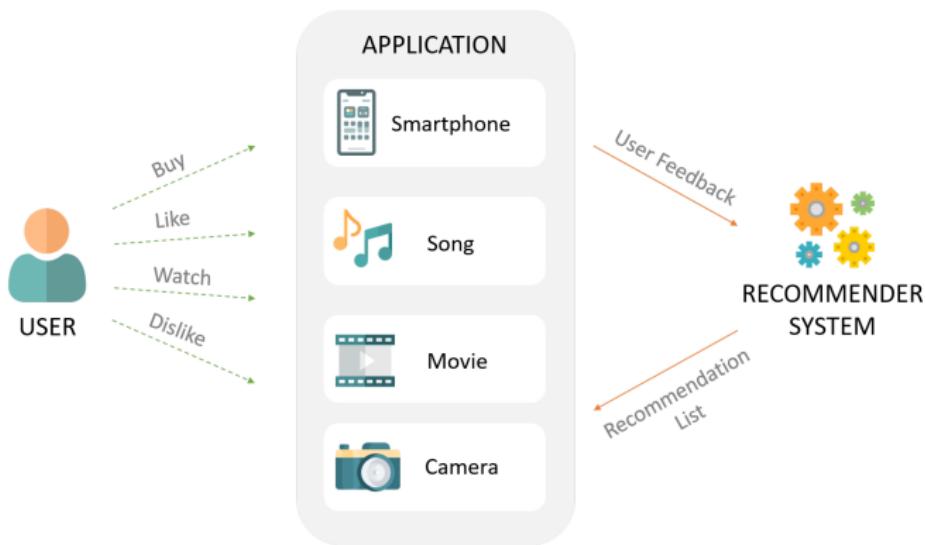
- **Gather** “known” ratings (populate matrix U)
- **Extrapolate** unknown ratings from known ones (mainly interested in high ratings, Top k)
 - For each user c or a subset of users, find

$$R_c = \arg \max_p U(c, p)$$

- R_c is the recommendation(s) for user c
- **Evaluate** extrapolation methods

Recommendation Systems: Challenges

- Explicitly survey users
- Implicitly learn ratings, e.g. purchase/suggestion to friend implies high rating
- Cold-start problem (new user, new product)



Recommendation Systems: Evaluation

	p_1	p_2	p_3	p_j								p_m
u_1	1	2	1	4			2	3	2	5		2
u_2		1			2		1		2		1	
u_3	1	1	2				1				1	2
			3	2			5		2		3	4
	1		2							5		
u_i	3	2	1	4	5		1	3	1	2		1
	4										4	
	5		1								5	
	1	4						1	3	5	1	2
u_n		3	1	1	2	1				4		5

- Compare predictions $U'(i,j)$ with known (hidden) ratings
- Root-mean-squared-error $\text{RMSE} = \sqrt{\sum_{i,j \in \text{Test Set}} (U(i,j) - U'(i,j))^2 / |\text{Test Set}|}$
- Spearman's Rank correlation, Kindell's Tau

Spearman's rank correlation coefficient

- A measure of similarity between two variables based on rank
- Correlation between ranks of values of the variables

$$\rho_{xy} = \frac{\text{cov}(rg_X, rg_Y)}{\sigma_{rg_X} \sigma_{rg_Y}}$$

X	3	3	3	4	4	4	5	5	5	0	0	1	1	2	2	2
Y	4	3	4	5	4	5	4	4	3	0	1	0	1	3	2	3
Z	1	0	1	2	1	2	1	1	0	3	4	3	4	0	5	0

Spearman's rank correlation coefficient

- A measure of similarity between two variables based on rank
- Correlation between ranks of values of the variables

$$\rho_{xy} = \frac{\text{cov}(rg_x, rg_y)}{\sigma_{rg_x} \sigma_{rg_y}}$$

X	3	3	3	4	4	4	5	5	5	0	0	1	1	2	2	2
Y	4	3	4	5	4	5	4	4	3	0	1	0	1	3	2	3
Z	1	0	1	2	1	2	1	1	0	3	4	3	4	0	5	0
rg _x	8	9	10	11	12	13	14	15	16	1	2	3	4	5	6	7
rg _y	10	6	11	15	12	16	13	14	7	1	3	2	4	8	5	9
rg _z	4	5	6	7	8	9	10	11	12	13	15	14	16	1	2	3

$$\rho_{XY} = 0.8 \quad \rho_{XZ} = -0.1 \quad \rho_{YZ} = -0.3$$

Other goodness measures for recommenders

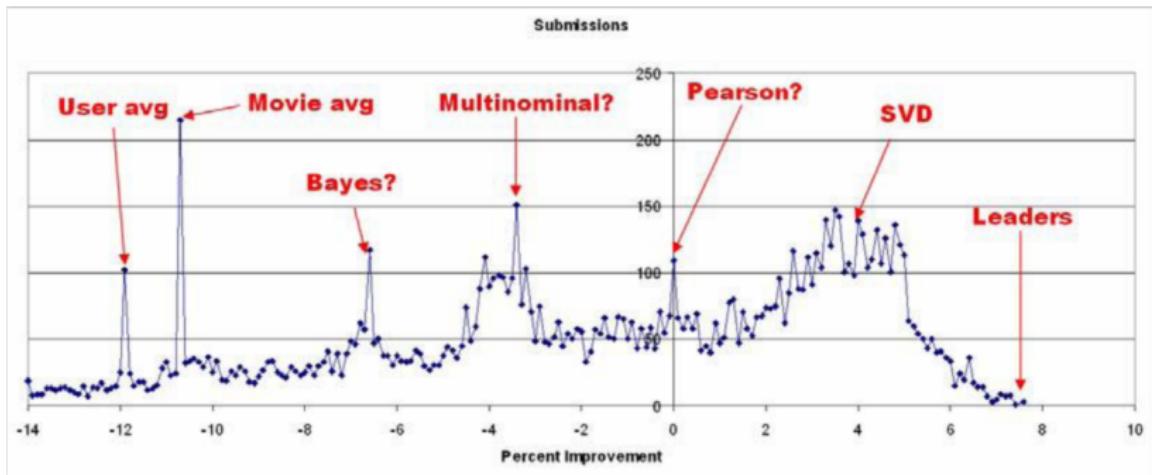
Other aspects of goodness of recommendations

- RMSE etc. might penalize methods for making errors on small rating
- **Prediction Diversity:** Customers are satisfied with intra-list diversity
- **Persistence:** re-show recommendations
- **Privacy:** User profiling and data gathering
- **Demographics:** important for customer satisfaction with recommendations
- **Trust:** Explaining how recommendations are found help
- **Serendipity:** How surprising recommendations are

Recommendation Systems: The Netflix Challenge

- In 2006, Netflix released a dataset movie rating dataset
- **Training set:** $\sim 1M$ ratings of the form $\langle \text{user}, \text{movie}, \text{date of grade}, \text{grade} \rangle$, 480,189 users, 17,770 movies
- **qualifying data set:** 2,817,131 triplets, $\langle \text{user}, \text{movie}, \text{date} \rangle$
- grade was known to jury only
- Competition to predict grades of qualifying set that improve accuracy by 10%
- Teams were informed of accuracy on 1,408,342 ratings (**validation set**)
- Jury used the test set of 1,408,789 ratings to determine winner

Netflix Challenge Methods

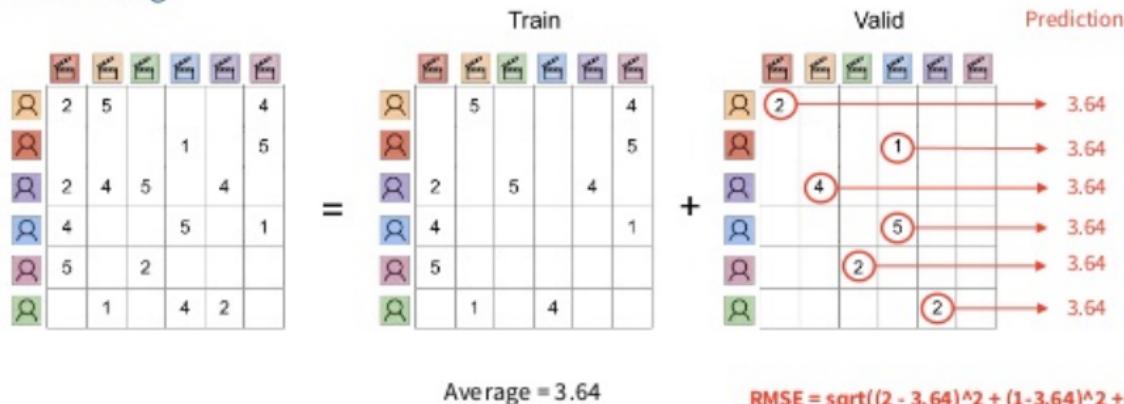


Recommendation Methods: Raw Averages

Predicting $U(i, j)$

- Assign average rating of all users for any item (MATRIX-AVERAGE)
- Assign average rating of all users for item j , (COLUMN-AVERAGE)
- Assign average rating of all items by user u , (ROW-AVERAGE)
- Mean is an unstable statistics (could use other measures of location)

Global Average



Recommendation Methods: Raw Averages

Predicting $U(i, j)$

Let MoM be the matrix mean (mean of means)

$$\text{MoM} := \frac{1}{nm} \sum_{c,p \in \text{Train Set}} U(c, p), \quad \text{then}$$

$$U'(i, j) = \text{MoM}$$

RMSE is just the standard-deviation of the data

Recommendation Methods: Raw Averages

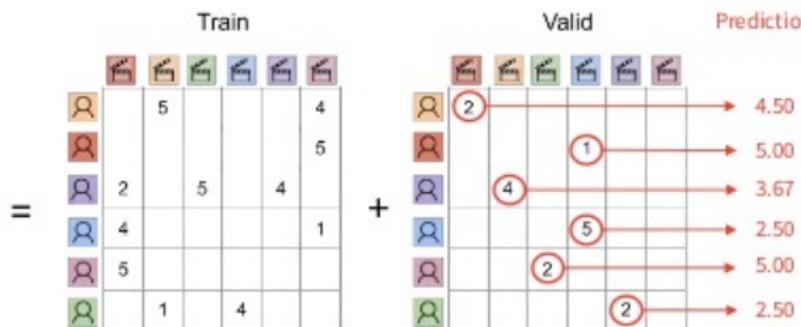
Predicting $U(i, j)$

- Assign average rating of all users for any item (MATRIX-AVERAGE)
- Assign average rating of all users for item j , (COLUMN-AVERAGE)
- Assign average rating of all items by user u , (ROW-AVERAGE)
- Mean is an unstable statistics (could use other measures of location)

User average

	Item 1	Item 2	Item 3	Item 4	Item 5	Item 6
User 1	2	5				4
User 2			1			5
User 3	2	4	5		4	
User 4			5			1
User 5	5		2			
User 6		1		4	2	

Recommender Systems from A to Z @ Crossing Minds



$$\text{Average } u1 = 4.50$$

$$\text{Average } u2 = 5.00$$

$$\text{Average } u3 = 3.67$$

$$\text{Average } u4 = 2.50$$

$$\text{Average } u5 = 5.00$$

$$\text{Average } u6 = 2.50$$

$$\text{RMSE} = \sqrt{(2 - 4.5)^2 + (1 - 5.0)^2 + \dots}$$

$$\text{RMSE} = \sqrt{6.15}$$

Recommendation Methods: ANOVA

Predicting $U(i, j)$

- Idea: Assign global average (matrix mean) $U(i, j) = \text{MoM}$
- Refinement 1: Product j maybe very (un) popular - highly (un)liked
 - Adjust for this bias
- Let dev_j be the average deviation of item j from MoM (+ve or -ve)
- $U(i, j) = \text{MoM} + dev_j$
- Refinement 2: User i may be very (non) pessimistic (critical)
 - Adjust for this bias too
- Let dev_i be the average deviation of user i from MoM

$$U(i, j) = \text{MoM} + dev_j + dev_i$$

Other methods are generally compared with this baseline

Recommendation Methods: Bayesian Method

Predicting $U(i, j)$

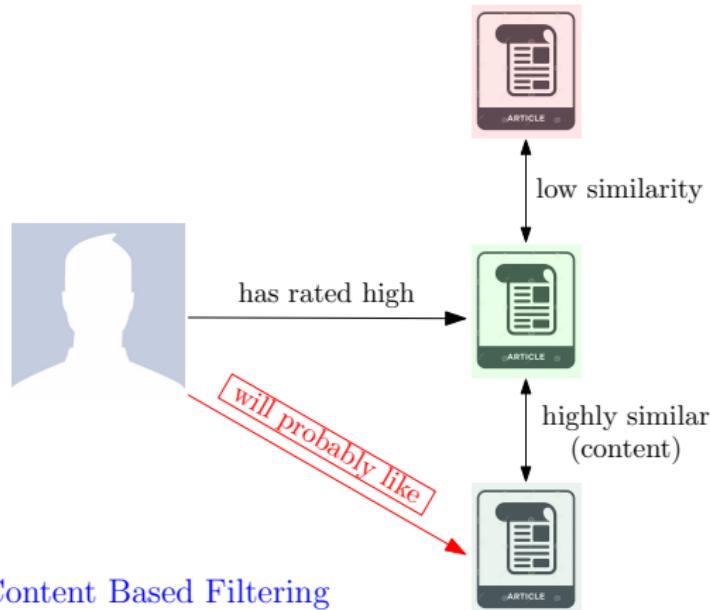
- Give the rating r , that is the most likely (in Bayesian sense)
- For $r \in \{0, 1, 2, 3, 4, 5\}$
- Let $P(r | j)$ be the conditional probability of rating r given item j

$$P(r | j) = \frac{P(j | r) P(r)}{P(j)}$$

- $P(j | r)$ is the conditional probability of item j given rating r
 - Estimate: The fraction of item j among all items rated with r
- $P(r)$ is the prior probability of rating r
- $P(j)$ is the prior probability of item j
- Doesn't take into account user i at all

Recommendation Methods: Content-based

Broad idea of Content Based Filtering



Recommendation Methods: Content-based

If j is similar to the “taste” of i , then predict $U(i,j)$ high

1 Build Item Profile (based on content) e.g.

- movies: vector of genre, director, budget, cast, plot, language
- books, blogs, website, news items: TF-IDF vector, author, topic

2 Build User Profile

- A vector with the same coordinates as item profile
- kind of “an average item” that the user likes **the taste of user**
- Weighted (by ratings) average of the item profiles that the user has rated

3 $U'(i,j)$: is \propto (cosine) similarity between item j 's and user i 's profiles

Recommendation Methods: Content-based

If j is similar to the “taste” of i , then predict $U(i,j)$ high

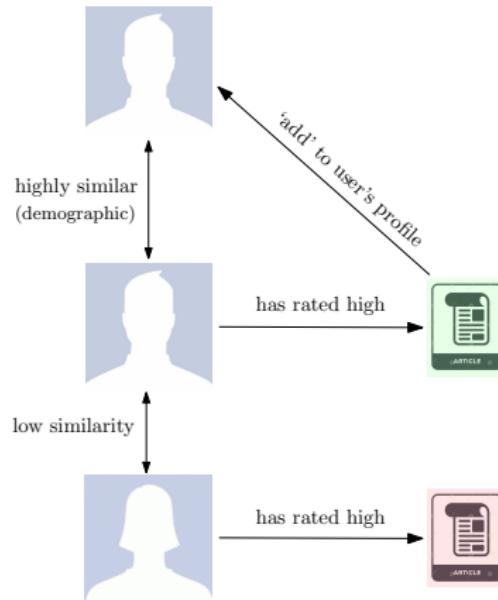
- Pros**
- No need of other users' information
 - No cold-start or sparsity problem w.r.t items
 - Unique taste of user is captured
 - Able to provide explanation (by listing contents features)
-
- Cons**
- Building profile is hard, find relevant features is hard
 - Cold start problem w.r.t users
 - User profile is heuristic
 - Overspecialization-never recommends items outside user profile
 - Does not cater for multiple interests of a user
 - Does not utilize judgment of other users

Recommendation Methods: Content-based

If j is similar to the “taste” of i , then predict $U(i,j)$ high

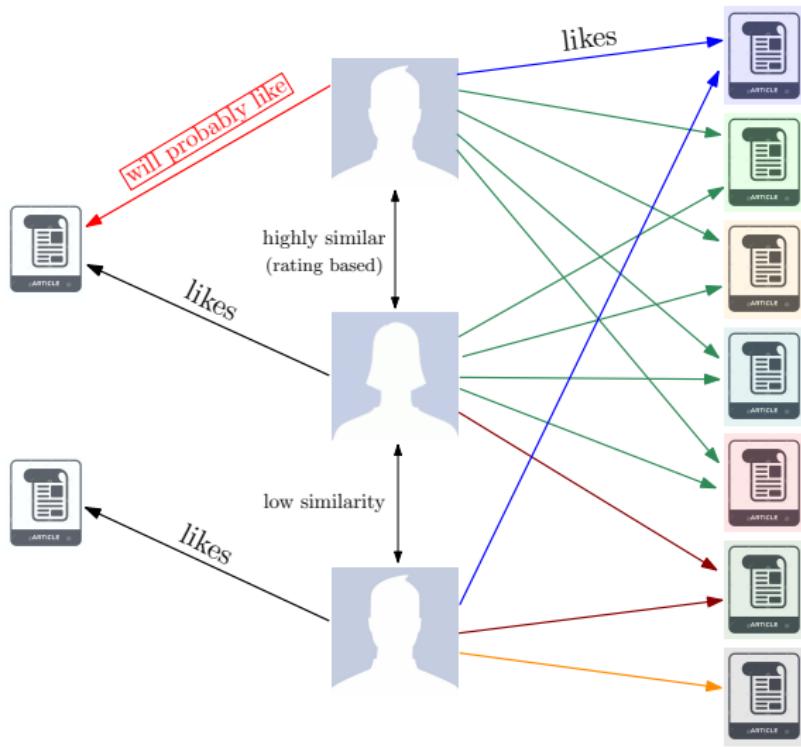
Can take into account other (similar) users judgments as follows.

Somewhat cater for the cold start problem



User-User Collaborative Filtering

Collaboratively filter (personalize) ratings using only the rating matrix U



User-User Collaborative Filtering

Collaboratively filter (personalize) ratings using only the rating matrix U

- Find the set N of users **with similar ratings** as of i
 - Find the top k similar rows to the i th row
- Estimate $U(i,j)$ as an “average” of $U(a,j)$ ’s for $a \in N$
- i has similar ‘taste’ to $a \in N \implies U(i,j)$ similar to $U(a,j)$

$$U'(i,j) = \frac{\sum_{a \in N} sim(a,i) \times U(a,j)}{\sum_{a \in N} sim(a,i)}$$

- Use cosine similarity to get N (\because interested in similar not high ratings)
- Alternatively, Spearman’s Rank Correlation, Kindell Tau

User-User Collaborative Filtering

Collaboratively filter (personalize) ratings using only the rating matrix U

- Find the set W of items **similarly rated** as j
 - Find the top k similar columns to the j th row
- Estimate $U(i, j)$ as an “average” of $U(i, p)$ ’s for $p \in W$

$$U'(i, j) = \frac{\sum_{p \in W} U(i, p) \times sim(j, p)}{\sum_{p \in W} sim(j, p)}$$

- Better result by item-item collaborative filtering
 - Because items are easier to model
 - has less complexity than users

Collaborative Filtering

- Works for any kind of items, unlike content based that requires outside knowledge for profiles
- **Cold-Start problem:** need enough users to find a match
- **Sparsity:** Hard to find users that have rated the same items
- **First rater:** cannot recommend items that are not previously rated (new or esoteric items)
- **Popularity bias:** Does not cater for unique taste of a user, tends to recommend popular items
- **Computational Complexity**

Generally, hybrid methods (ensemble) are used. Combine predictions of many recommender system (average, weighted average, regression)